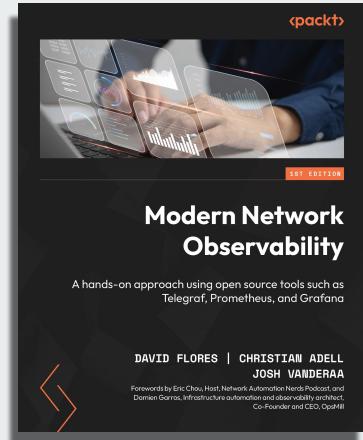

Modern Network Observability Workshop

CfgMgmtCamp (04/02/2026)
Ghent

David Flores & Christian Adell



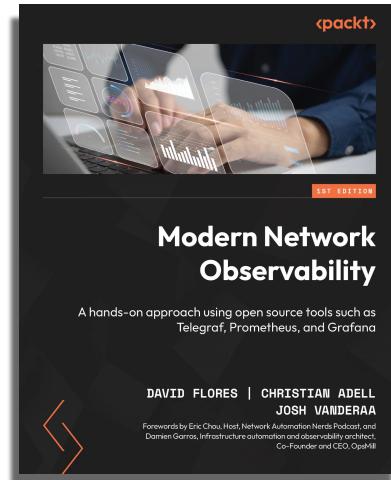
About Us: David Flores

- Network Automation and Observability at CoreWeave
- Experience in Service Provider, Financial, Hyperscalers, and Integrators.
- Focused on building open source automation and network observability platforms
- Co-Author of “Modern Network Observability”

[@netpanda.bsky.social](https://bluesky.net/@netpanda.bsky.social) - Bluesky

[@davidban77](https://github.com/davidban77) - Github / X

[David Flores](https://www.linkedin.com/in/david-flores-11111111) - LinkedIn



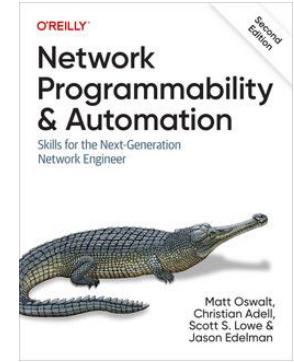
About Us: Christian Adell

- Network Automation and Observability at CoreWeave
- Broad Experience: Service and Content Provider, Enterprises, Hyperscalers, and Integrators.
- Co-Author of “Modern Network Observability” and other books
- Currently, writing an open book at

<https://designingnetworkautomation.com/>

[@chadell](#) - Github

[Christian Adell](#) - LinkedIn





Workshop Agenda

1. From Monitoring to Observability: Understanding the Framework
2. Data Collection & Normalization: Building Trustworthy Signals
3. Turning Data into Insight: Alerting, Queries & Visualization
4. Day-2 Automation: From Observability to Intelligent Workflows

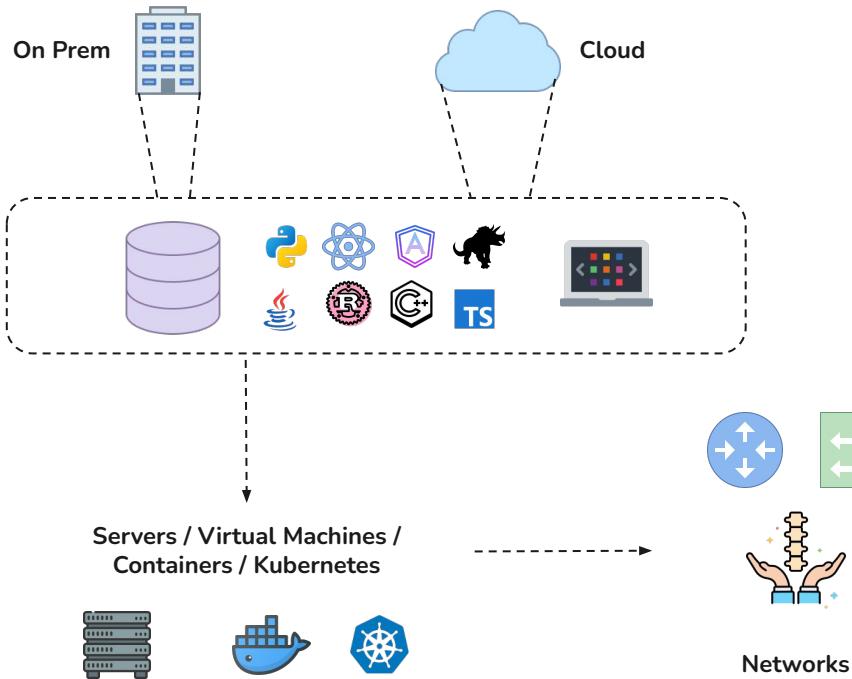
All the code is at: <https://github.com/network-observability/network-observability-lab>

- Base scenario: **webinar**
- Scenario solved (spoiler alert): **webinar-completed**

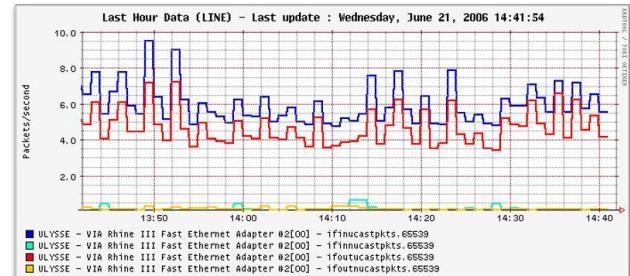
1

From Monitoring to Observability: Understanding the Framework

Network Monitoring Today



- MRTG, PRTG, Cacti – Familiar but limited.
- **Networks: The backbone of modern infrastructure.**
- Complexity is growing: cloud, on-premises, hybrid, microservices.
- Challenge? some metrics and disparate logging isn't enough for deep insights.



What Is Observability?

- **Observability vs. Monitoring:** “What happened” vs. “Why it happened.”
- Extracting actionable insights from **multiple system signals**.
- Distributed systems and modular architectures require new strategies.

Synthetic Monitoring

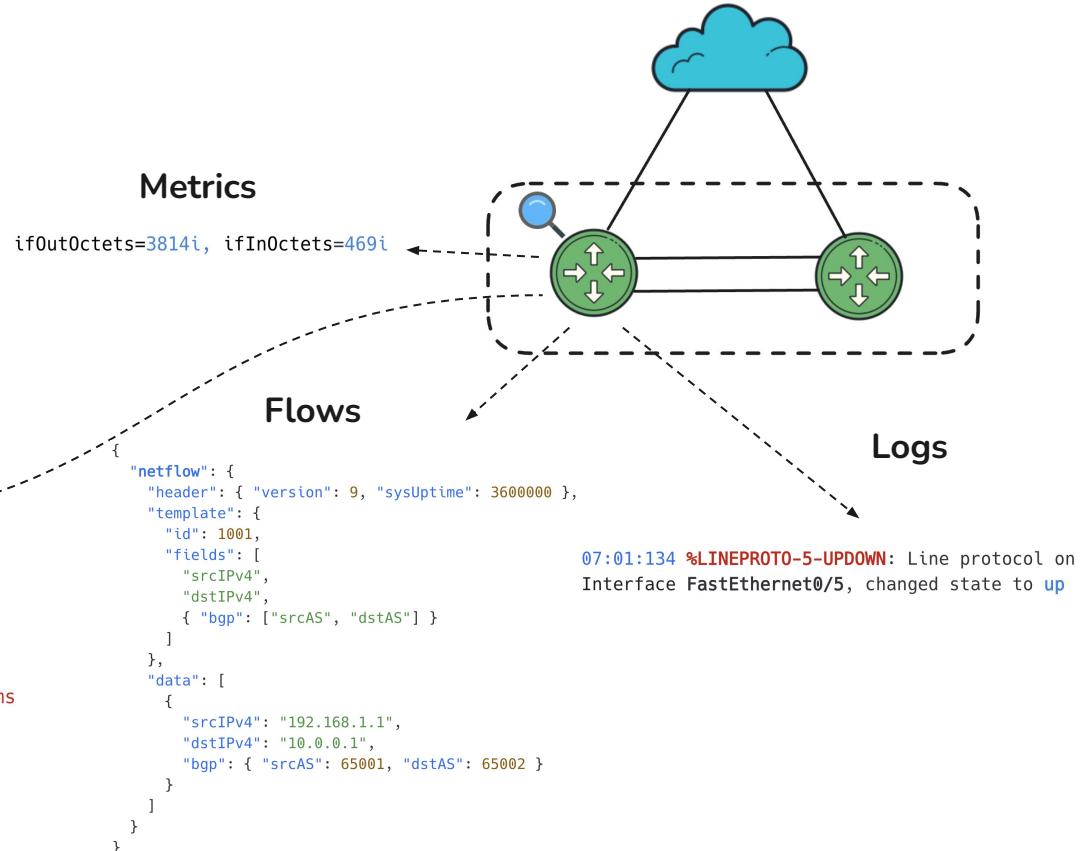
```
Router# ping 192.0.2.1
```

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 192.0.2.1, timeout is 2 seconds:

!!!!

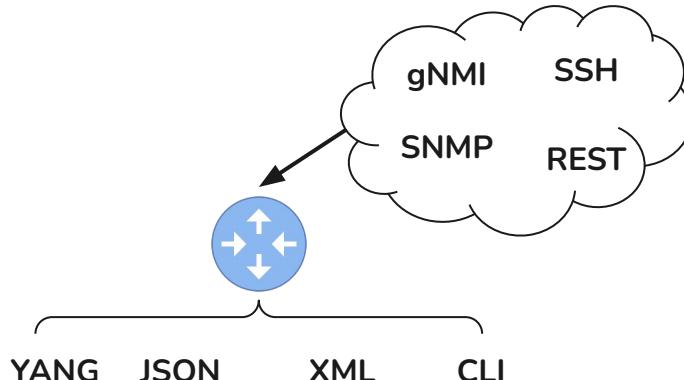
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/6/8 ms



The Problem with Current Tools



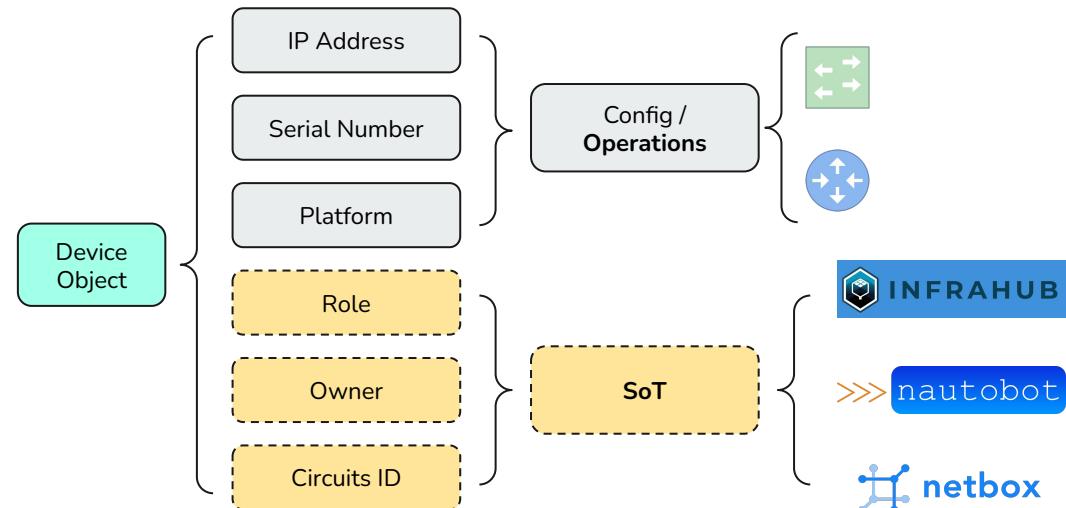
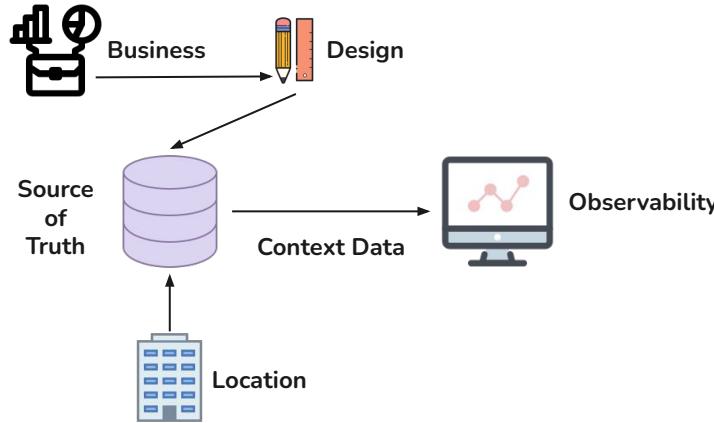
- **Vendor Lock-In:** Limited scalability and flexibility.
- **Static Dashboards:** Outdated visualizations that don't adapt.
- **Inflexible Data Collection and Processing:** Data comes in many many many forms...



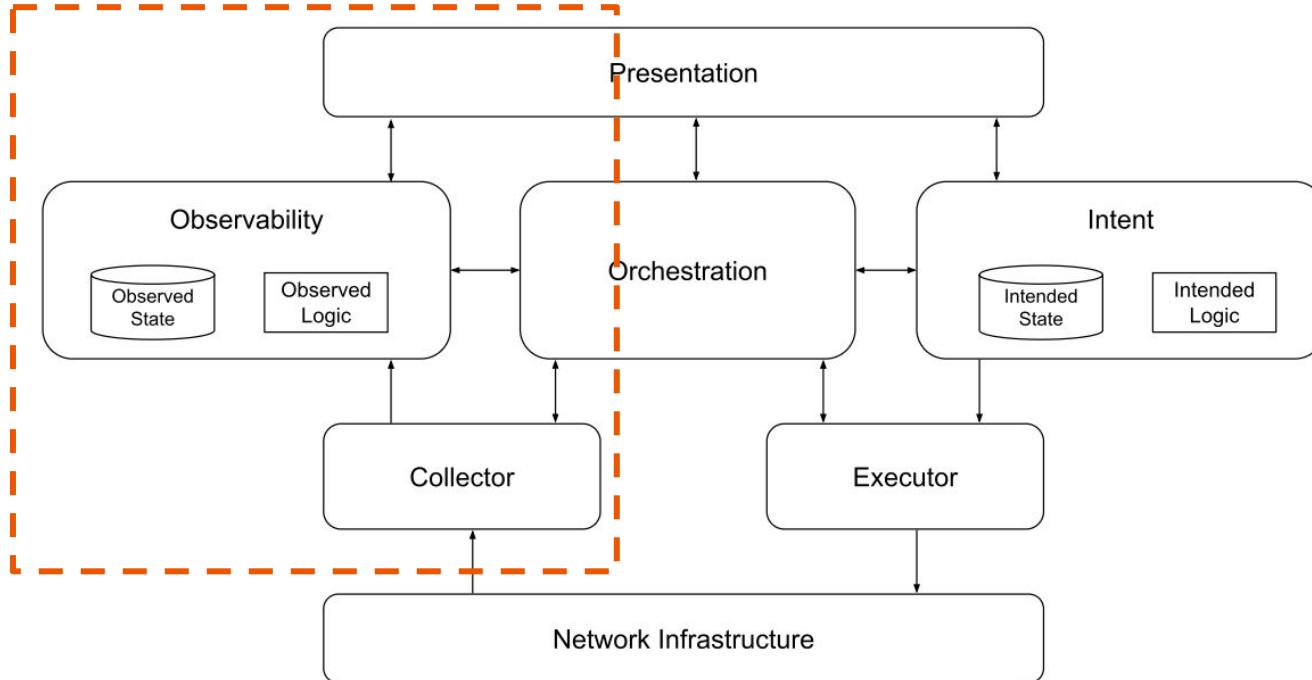
The Problem with Current Tools



- **Missing Context:** Lack of contextual (enrichment) data for smarter alerts.



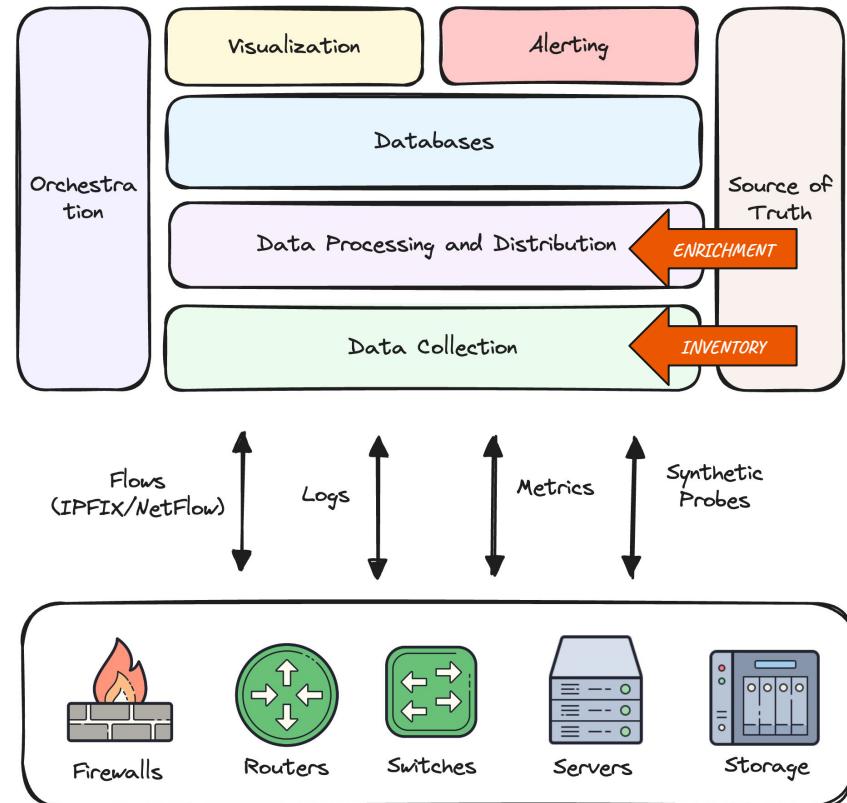
Let's talk about Architecture



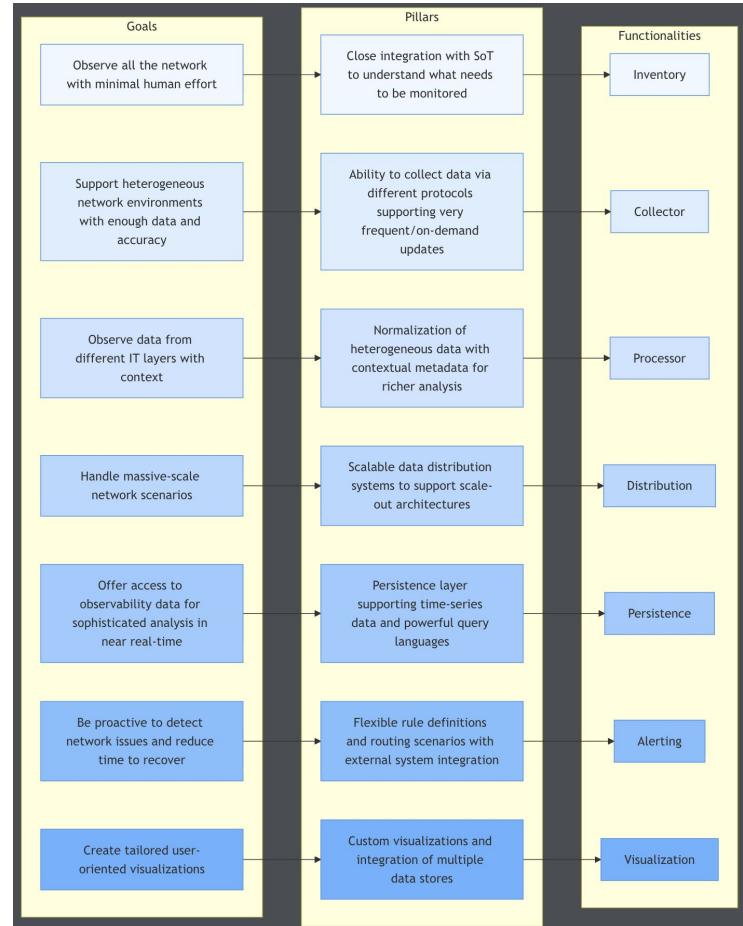
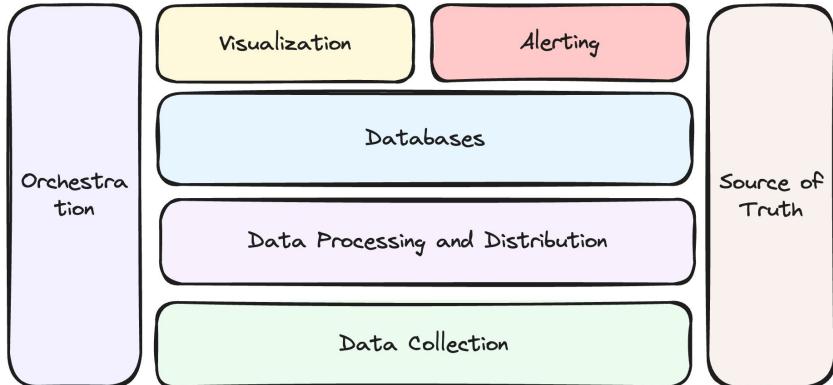
The Observability Architecture

Modular architecture and specialized components help address the challenges.

- **Collection:** Gather data from SNMP, streaming telemetry, logs.
- **Processing & Distribution:** Clean and unify diverse data.
- **Storage:** Leverage specialized databases (time-series, log storage).
- **Visualization & Alerts:** Turn data into actionable insights with tools like Grafana.



Goals - Pillars - Components

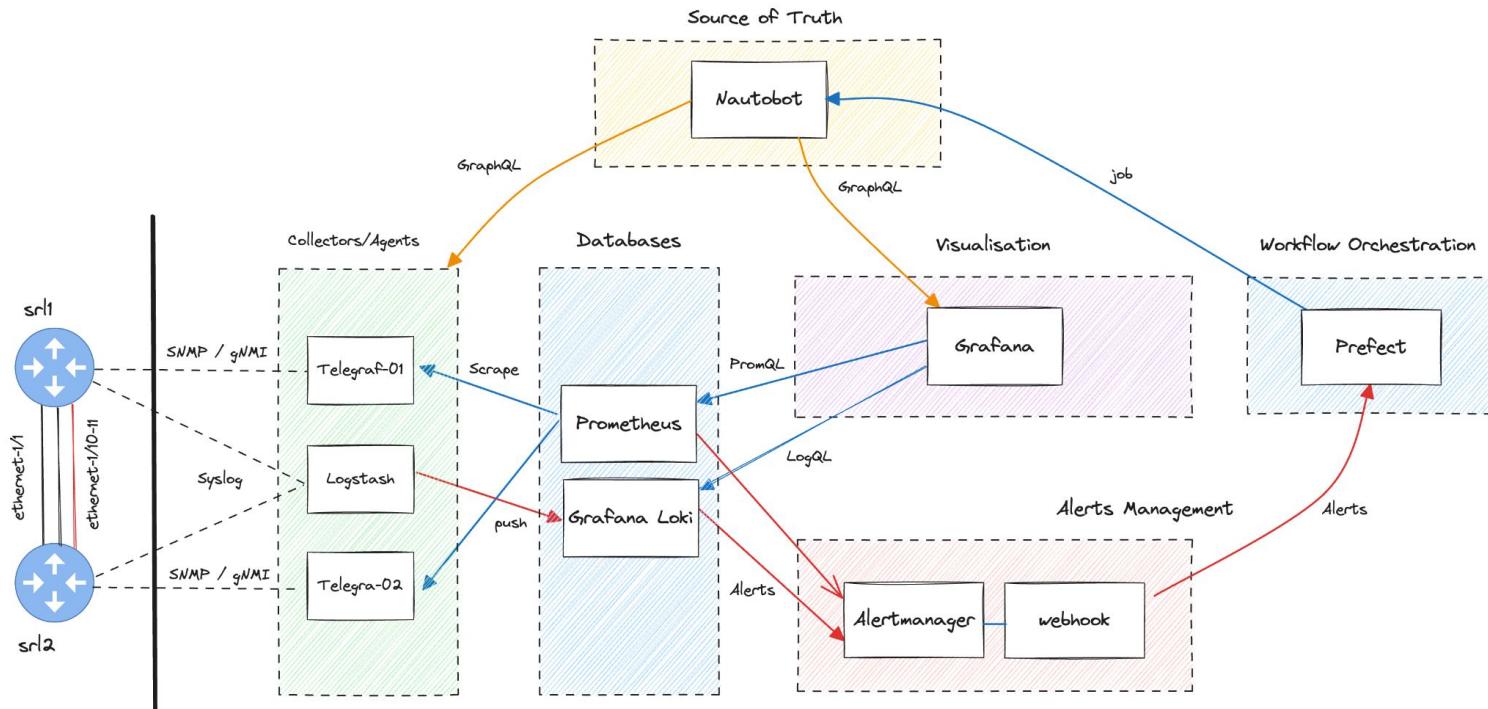


Exercise 1: Setting Up the Stack

Workshop Exercises

1. Setting Up the Stack
2. Build Your First Normalized Telemetry Pipeline
3. Build Alerts & Explore Data
4. Build an Observability-Driven Automation Flow

Workshop Scenario



Disclaimer: None of the elements of the scenario should be seen as a recommendation. They are used as examples, and are completely interchangeable by the tools you prefer

Setup Options

All the code: <https://github.com/network-observability/network-observability-lab>

Two main options:

1. Setup in Linux with Ansible
2. Unattended install in Digital Ocean

Setup in Linux with Ansible (1/2)

1. Fork (in GitHub) and Clone the Git Repository

```
$ git clone https://github.com/<your-user>/network-observability-lab.git  
$ cd network-observability-lab
```

2. Install the netobs tool (use a virtualenv!)

```
(.venv)$ pip install .
```

3. Setup environment files

```
$ cp example.env .env
```

Setup in Linux with Ansible (2/2)

4. Deploy the stack with Docker (dependency)

```
(.venv)$ netobs lab deploy \
    --scenario webinar \
    --topology ./chapters/webinar/containerlab/lab.yml \
    --vars-topology ./chapters/webinar/containerlab/lab_vars.yml
```



Unattended install in Digital Ocean (1/3)

1. Create a Digital Ocean account: <https://www.digitalocean.com/try/free-trial-offer>
2. Fork (in GitHub) and Clone the Git Repository

```
$ git clone https://github.com/<your-user>/network-observability-lab.git  
$ cd network-observability-lab
```

3. Install the netobs tool (use a virtualenv!)

```
(.venv)$ pip install .
```



Unattended install in Digital Ocean (2/3)

4. Create a Digital Ocean API token

5. Create an SSH key pair and retrieve its fingerprint

```
$ ssh-keygen -t rsa -b 4096 -C "network-observability-lab" -f ~/.ssh/id_rsa_do
```

6. Setup environment files

```
$ cp example.setup.env .setup.env  
  
$ cp example.env .env  
  
SSH_KEY_PATH=~/.ssh/id_rsa_do  
  
NETOBS_REPO="https://github.com/<your-username>/network-observability-lab.git"
```

Unattended install in Digital Ocean (3/3)

7. Create a Digital Ocean droplet (it takes some minutes)

```
(.venv) $ netobs setup deploy --scenario webinar \
    --topology ./chapters/webinar/containerlab/lab.yml \
    --vars-topology ./chapters/webinar/containerlab/lab_vars.yml
[17:15:12] Running command: ansible-playbook setup/create_droplet.yml -i
setup/inventory/localhost.yaml
Enter the droplet image [ubuntu-22-04-x64]:
Enter the droplet size [s-8vcpu-16gb]: s-4vcpu-8gb
Enter the droplet region [fra1]:
PLAY [Stand up netobs-droplet]
*****
```



Verifying the Stack (1/4)

- ❑ (if using DO) Connect to remote server

```
$ netobs setup show
[08:45:17] Running command: ansible-playbook setup/list_droplet.yml -i
setup/inventory/do hosts.yaml
PLAY [Show Inventory] ****
TASK [Show SSH command] ****
ok: [netobs-droplet] => {}
MSG:
ssh -o StrictHostKeyChecking=no -i ~/.ssh/id_rsa_do root@192.0.2.1
```

- ❑ Connect via SSH

```
$ ssh -o StrictHostKeyChecking=no -i ~/.ssh/id_rsa_do root@192.0.2.1
```

Verifying the Stack (2/4)

```
root@netobs-droplet:~/network-observability-lab# netobs lab show --scenario webinar --topology chapters/webinar/containerlab/lab.yml
[17:57:23] Showing lab environment for scenario: webinar
  Showing containers for service(s): []
    Running command: docker compose --project-name netobs -f chapters/webinar/docker-compose.yml --verbose ps
      NAME          IMAGE           COMMAND          SERVICE          CREATED         STATUS          PORTS
      alertmanager  docker.io/prom/alertmanager:v0.26.0  "/bin/alertmanager -..."  alertmanager    39 minutes ago Up 39 minutes  0.0.0.0:9093->9093/tcp,
      [:]:9093/tcp
      grafana       docker.io/grafana/grafana:10.4.4   "/run.sh"        grafana        39 minutes ago Up 39 minutes  0.0.0.0:3000->3000/tcp,
      [:]:3000->3000/tcp
      logstash      docker.io/grafana/logstash-output-loki:3.1.1  "/usr/local/bin/dock..."  logstash       39 minutes ago Up 39 minutes  0.0.0.0:1515->1515/tcp,
      [:]:1515->1515/tcp, 0.0.0.0:9600->9600/tcp, [:]:9600->9600/tcp, 5044/tcp, 0.0.0.0:12201->12201/udp, [:]:12201->12201/udp
      loki          docker.io/grafana/loki:3.1.1     "/usr/bin/loki -conf..."  loki          39 minutes ago Up 39 minutes  0.0.0.0:3001->3001/tcp,
      [:]:3001->3001/tcp, 3100/tcp
      nautobot     docker.io/networktocode/nautobot:2.4-py3.10  "/dockerc-entrypoint..."  nautobot      39 minutes ago Up 39 minutes (unhealthy)  0.0.0.0:8080->8080/tcp,
      [:]:8080->8080/tcp, 0.0.0.0:8443->8443/tcp, [:]:8443->8443/tcp
      nautobot-postgres  postgres:14          "dockerc-entrypoint.s..."  nautobot-postgres  39 minutes ago Up 39 minutes  5432/tcp
      nautobot-redis   redis:7.2-alpine      "dockerc-entrypoint.s..."  nautobot-redis   39 minutes ago Up 39 minutes  6379/tcp
      netobs-postgres-1  postgres:14          "dockerc-entrypoint.s..."  postgres        39 minutes ago Up 39 minutes (healthy)  5432/tcp
      netobs-prefect-server-1  prefecthq/prefect:3-latest  "/usr/bin/tini -g ..."  prefect-server  39 minutes ago Up 39 minutes  0.0.0.0:4200->4200/tcp,
      [:]:4200->4200/tcp
      netobs-prefect-services-1  prefecthq/prefect:3-latest  "/usr/bin/tini -g ..."  prefect-services  39 minutes ago Up 39 minutes  8092/udp, 8125/udp, 8094
      netobs-redis-1    redis:7            "dockerc-entrypoint.s..."  redis          39 minutes ago Up 39 minutes (healthy)  6379/tcp
      prometheus      docker.io/prom/prometheus:v2.52.0   "/bin/prometheus -s..."  prometheus     39 minutes ago Up 39 minutes  0.0.0.0:9090->9090/tcp,
      [:]:9090->9090/tcp
      telegraf-01     netobs-telegraf-01    "/entrypoint.sh tele..."  telegraf-01     39 minutes ago Up 39 minutes  8092/udp, 8125/udp, 8094
      /tcp, 0.0.0.0:9004->9004/tcp, [:]:9004->9004/tcp
      telegraf-02     netobs-telegraf-02    "/entrypoint.sh tele..."  telegraf-02     39 minutes ago Up 39 minutes  8092/udp, 8125/udp, 8094
      /tcp, 0.0.0.0:9005->9005/tcp, [:]:9005->9005/tcp
      webhook        netobs-webhook      "python -m app.main"    webhook        39 minutes ago Up 39 minutes  0.0.0.0:9997->9997/tcp,
      [:]:9997->9997/tcp
      Successfully ran: show containers
      End of task: show containers
```

```
Showing containerlab topology
Topology file: chapters/webinar/containerlab/lab.yml
Running command: sudo containerlab inspect -t chapters/webinar/containerlab/lab.yml
17:57:23 INFO Parsing & checking topology file:lab.yml
```

Name	Kind/Image	State	IPv4/6 Address
clab-mno-webinar-srl1	nokia_srlinux ghcr.io/nokia/srlinux	running	198.51.100.3 N/A
clab-mno-webinar-srl2	nokia_srlinux ghcr.io/nokia/srlinux	running	198.51.100.2 N/A

```
Successfully ran: Inspect containerlab topology
```



Verifying the Stack (3/4)

- ❑ Access SRLinux boxes and validate connectivity over virtual link (creds: *admin/NokiaSrl1!*)

```
root@netobs-droplet:~/network-observability-lab# ssh admin@srl1
Warning: Permanently added 'srl1' (ED25519) to the list of known hosts.
(admin@srl1) Password:
Last login: Tue Jan  6 11:58:42 2026 from 198.51.100.1
Loading environment configuration file(s): ['/etc/opt/srlinux/srlinux.rc']
Welcome to the Nokia SR Linux CLI.
--{ + running }--[ ]--
A:admin@srl1# ping 10.1.2.2 network-instance default
Using network instance default
PING 10.1.2.2 (10.1.2.2) 56(84) bytes of data.
64 bytes from 10.1.2.2: icmp seq=1 ttl=64 time=115 ms
64 bytes from 10.1.2.2: icmp_seq=2 ttl=64 time=4.24 ms
```



Verifying the Stack (4/4)

- ❑ Access **Grafana** at <http://<your-ip>:3000>
- ❑ Access **Nautobot** at <http://<your-ip>:8080>

The screenshot shows the Grafana home page. The left sidebar has links for Home, Starred, Dashboards, and Alerting. The main area displays a dashboard titled "CAN data analysis with Grafana Assistant". This dashboard includes a map and several data cards showing metrics like 100s, 88s, 3.50, 14.2 ms, and 98.0 km. Below the dashboard, there's a link to a blog post: <https://grafana.com/blog/how-to-use-ai-to-analyze-and-visualize-can-data-with-grafana-assistant/>.

The screenshot shows the Nautobot login page. The URL is <http://161.35.73.188:8080/login/?next=/>. The page features a logo with three orange arrows pointing right, followed by the text "nautobot". A "Log in" button is visible at the bottom right of the login form.



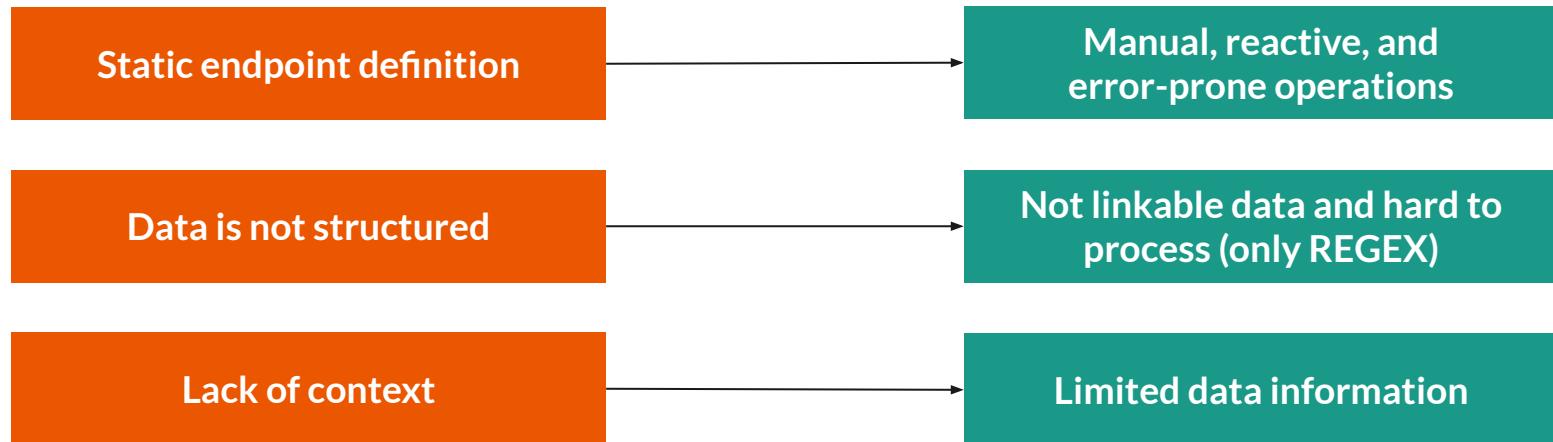
Wrap Up

- Observability goes beyond static monitoring to understand the WHY
- Within a complete network automation architecture, Observability is a key enabler
- Understanding the functionalities enables ownership

2

Data Collection & Normalization: Building Trustworthy Signals

Challenges of Traditional Data Collection

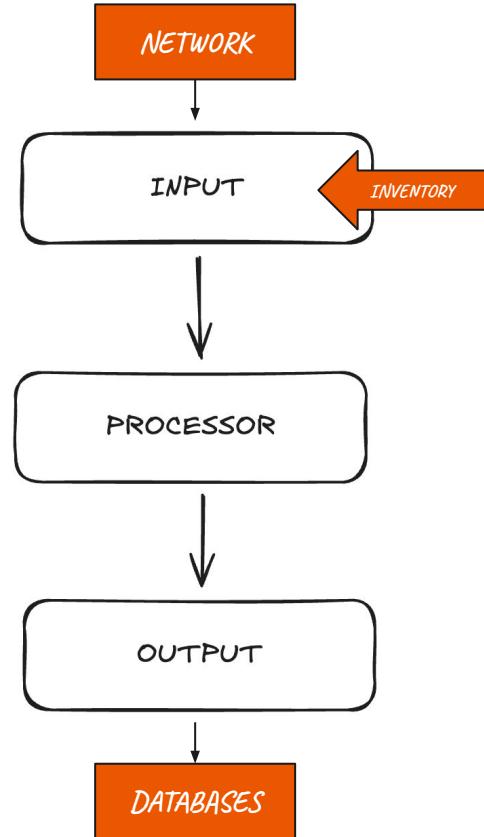


What Data Collection should look like?

- The target endpoints should be **dynamically** obtained and integrated with
- All the data type are **structured** according to a consistent data structure
- The data structure should incorporate **metadata** (labels) that allow connecting data
- The data identifier, value and labels should be able to adapt align with **standardization**
- The collector layer should be flexible to adapt to **different types** of data (using multiple tools if needed)
- The collectors should support **scale out** architecture to adapt to how the network infrastructure evolves

Collectors' Architecture

- **Input:** Defines what to collect (pull or push), the target and the paths
- **Processor:** Defines how to transforms the data: normalization and enrichment
- **Output:** Defines how to expose (passive mode) or export the data (active mode)





Classifying Data Inputs

Data Type	Protocols / Collection Methods	Notes / Examples
Metrics	SNMP, HTTP scraping, CLI polling, OpenTelemetry (OTLP), Streaming telemetry (gNMI)	Device metrics, host metrics, application metrics
Logs	OpenTelemetry (OTLP), file tailing, syslog	Application logs, system logs, structured logs
Traces	OpenTelemetry (OTLP)	Distributed tracing across services
Network Flows	NetFlow, IPFIX	Traffic flows, source/destination analysis
Protocol-specific	BMP, BGP, ARP, OSPF	BGP monitoring (BMP), ARP tables, BGP tables, OSPF tables
Packet Captures	PCAP (libpcap), SPAN / TAP	Full packet inspection, deep troubleshooting

Data Normalization

```
interface, collection_type=gnmi, host=telegraf-01, name=ethernet-1/53, source=srl1, admin_state="disable", oper_state="up" 1767723255964010964
```

```
interface, collection_type=snmp, host=telegraf-01, agent_host=srl1, ifName=ethernet-1/53, ifIndex=101 ifAdminStatus=1, ifOperStatus=2i 1767723255964010964
```

interface_oper_state {device="srl1", name="ethernet-1/1", intf_role="peer"} 2

Metric Name
Labels
Value

Metric Name Normalized

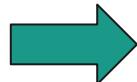
```
interface_oper_state {  
    name = "ethernet-1/53",  
    device = "srl1",  
    origin = "telegraf-01",  
} 1 @1767723255964010964
```

Labels Normalized:
- name
- device

Value Normalized:
- down: 0
- up: 1

Data Enrichment

```
interface_oper_state {  
    name = "ethernet-1/53",  
    device = "srl1",  
} 1@1767723255964010964
```



```
interface_oper_state {  
    name = "ethernet-1/53",  
    device = "srl1",  
    role = "uplink",  
    device_role = "border-router",  
    status = "maintenance",  
    location = "lab-1",  
    region = "Europe",  
} 1@1767723255964010964
```

Expanded Context



Data Output Options

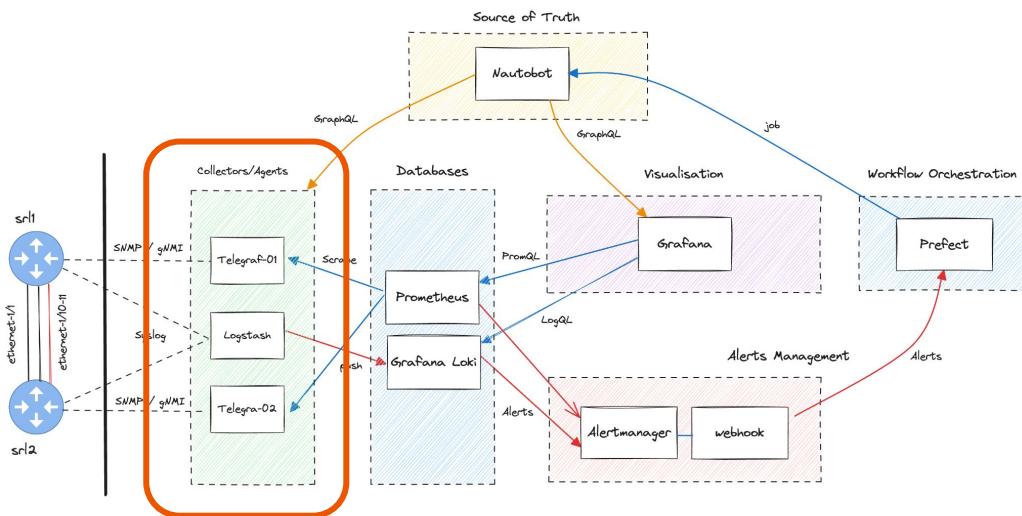
- Write locally (e.g., file or local database)
- Write remotely (e.g., remote writes)
- Streaming to a message bus (e.g., Kafka)
- Expose API endpoint (e.g., HTTP scraping)

Exercise 2.1: Build Your First Normalized Telemetry Pipeline Metrics with Telegraf

Workshop Exercises

1. Setting Up the Stack
2. **Build Your First Normalized Telemetry Pipeline**
3. Build Alerts & Explore Data
4. Build an Observability-Driven Automation Flow

Workshop Scenario



```
.
├── containerlab
│   ├── lab_vars.yml
│   └── lab.yml
└── startups
    ├── startup-config-sr1.json
    └── startup-config-sr2.json
├── docker-compose.yml
├── grafana
│   ...
└── logstash
    └── logstash.cfg
├── loki
│   ...
└── prometheus
│   ...
├── README.md
└── telegraf
    ├── telegraf-01.conf.toml
    ├── telegraf-02.conf.toml
    └── telegraf.Dockerfile
```

Telegraf Inputs: ICMP

```
[[inputs.ping]]
  interval = "10s"
  count = 3
  ping_interval = 1.0
  timeout = 5.0
  urls = ["srl1"]

[inputs.ping.tags]
  collection_type = "ping"
  watcher_type = "healthcheck"
  service = "srl1"
```



Telegraf Inputs: SNMP

```
[[inputs.snmp]]
agents = ["srl1"]
version = 2
community = "${SNMP_COMMUNITY}"
interval = "60s"
timeout = "10s"
retries = 3

[[inputs.snmp.tags]]
collection_type = "snmp"

[[inputs.snmp.table]]
name = "interface"

[[inputs.snmp.table.field]]
name = "name"
# IF-MIB::ifDescr
oid = "1.3.6.1.2.1.2.2.1.2"
is_tag = true
```

Telegraf Inputs: gNMI

```
[[inputs.gnmi]]
addresses = ["srl1:57400"]
username = "${NETWORK_AGENT_USER}"
password = "${NETWORK_AGENT_PASSWORD}"
redial = "20s"
tagexclude = ["path"]
tls_enable = true
insecure_skip_verify = true

[inputs.gnmi.tags]
collection_type = "gnmi"

[[inputs.gnmi.subscription]]
name = "interface"
path = "/interface[name=*/statistics"
subscription_mode = "sample"
sample_interval = "10s"
```

Other options:

- Kubernetes
- Docker
- Cloud provider services
- HTTP
- Databases
- Message Queues
- Tail (files, logs)
- EXEC



Telegraf Outputs

```
[[outputs.file]]
files = ["stdout"]

[[outputs.prometheus_client]]
# HTTP port to listen on
listen = ":9004"
```

Other options:

- Remote Write to TSDB:
Prometheus, VictoriaMetrics,
InfluxDB
- Message buses: Kafka, NATS, AWS
Kinesis
- Managed Services: Datadog,
NewRelic
- Files



Applying Telegraf changes

```
root@netobs-droplet:~/network-observability-lab# netobs lab update telegraf-01 -s webinar
[18:07:43] Updating lab environment for scenario: webinar
      Removing service(s): ['telegraf-01']
      Running command: docker compose --project-name netobs -f chapters/webinar/docker-compose.yml --verbose rm --stop --force --volumes telegraf-01
[+] stop 1/1
[+] remove 1/1telegraf-01 Stopped                               0.5s
✓ Container telegraf-01 Removed                                0.7s
[18:07:44] Successfully ran: remove containers
```

```
      Starting service(s): ['telegraf-01']
      Running command: docker compose --project-name netobs -f chapters/webinar/docker-compose.yml --verbose up -d --remove-orphans telegraf-01
WARN[0000] No services to build
[+] up 1/1
✓ Container telegraf-01 Created                                0.1s
[18:07:46] Successfully ran: start stack
```

Lab environment updated for scenario: webinar



Getting Telegraf Outputs (without processing)

```
root@netobs-droplet:~/network-observability-lab# netobs docker logs telegraf-01 --tail 10
[18:11:01] Showing logs for service(s): ['telegraf-01']

Running command: docker compose --project-name netobs -f chapters/batteries-included/docker-compose.yml logs --tail=10 telegraf-01
telegraf-01 | interface,agent_host=srl1,collection_type=snmp,host=telegraf-01,name=ethernet-1/40,in_octets=0i,out_octets=0i 1767723300000000000000
telegraf-01 | interface,agent_host=srl1,collection_type=snmp,host=telegraf-01,name=ethernet-1/3,in_octets=0i,out_octets=0i 1767723300000000000000
telegraf-01 | interface,collection_type=gnmi,host=telegraf-01,name=ethernet-1/53,source=srl1 admin_state="disable",oper_state="down" 1767723255964010964
telegraf-01 | interface,collection_type=gnmi,host=telegraf-01,name=ethernet-1/54,source=srl1 admin_state="disable",oper_state="down" 1767723255964010964
telegraf-01 | interface,collection_type=gnmi,host=telegraf-01,name=ethernet-1/55,source=srl1 admin_state="disable",oper_state="down" 1767723255964010964
telegraf-01 | interface,collection_type=gnmi,host=telegraf-01,name=ethernet-1/56,source=srl1 admin_state="disable",oper_state="down" 1767723255964010964
telegraf-01 | interface,collection_type=gnmi,host=telegraf-01,name=ethernet-1/57,source=srl1 admin_state="disable",oper_state="down" 1767723255964139249
telegraf-01 | interface,collection_type=gnmi,host=telegraf-01,name=ethernet-1/58,source=srl1 admin_state="disable",oper_state="down" 1767723255964139249
telegraf-01 | interface,collection_type=gnmi,host=telegraf-01,name=mgmt0,source=srl1 admin_state="enable",oper_state="up" 1767723255964139249
telegraf-01 | ping,collection_type=ping,host=telegraf-01,service=srl1,url=srl1,watcher_type=healthcheck
percent_packet_loss=0,minimum_response_ms=2.385,average_response_ms=3.747,result_code=0i,packets_transmitted=3i,packets_received=3i,ttl=64i,maximum_response_ms=4.677,standard_deviation_ms=0.984 1767723262000000000
```

Successfully ran: show logs

End of task: show logs

Implementing Telegraf Normalization

```
[[processors.rename]]  
[[processors.rename.tagpass]]  
collection_type = ["snmp"]
```

```
[[processors.rename]]  
[[processors.rename.tagpass]]  
collection_type = ["gnmi"]
```

```
[[processors.rename]]  
[[processors.rename.tagpass]]  
collection_type = ["ping"]
```

```
[[processors.rename.replace]]  
tag = "agent_host"  
dest = "device"
```

```
[[processors.rename.replace]]  
tag = "source"  
dest = "device"
```

```
[[processors.rename.replace]]  
tag = "url"  
dest = "device"
```



Implementing Telegraf Enrichment

```
[ [processors.regex]]
namepass = ["interface"]
[processors.regex.tagpass]
device = ["srl1"]
[ [processors.regex.tags]]
key = "name"
pattern = "^Ethernet.*$"
result_key = "intf_role"
replacement = "peer"
[ [processors.regex.tags]]
key = "name"
pattern = "^Management.*$"
result_key = "intf_role"
replacement = "mgmt"
```

Other options:

- override
- converter
- enum
- filter
- starlark
- printer
- date



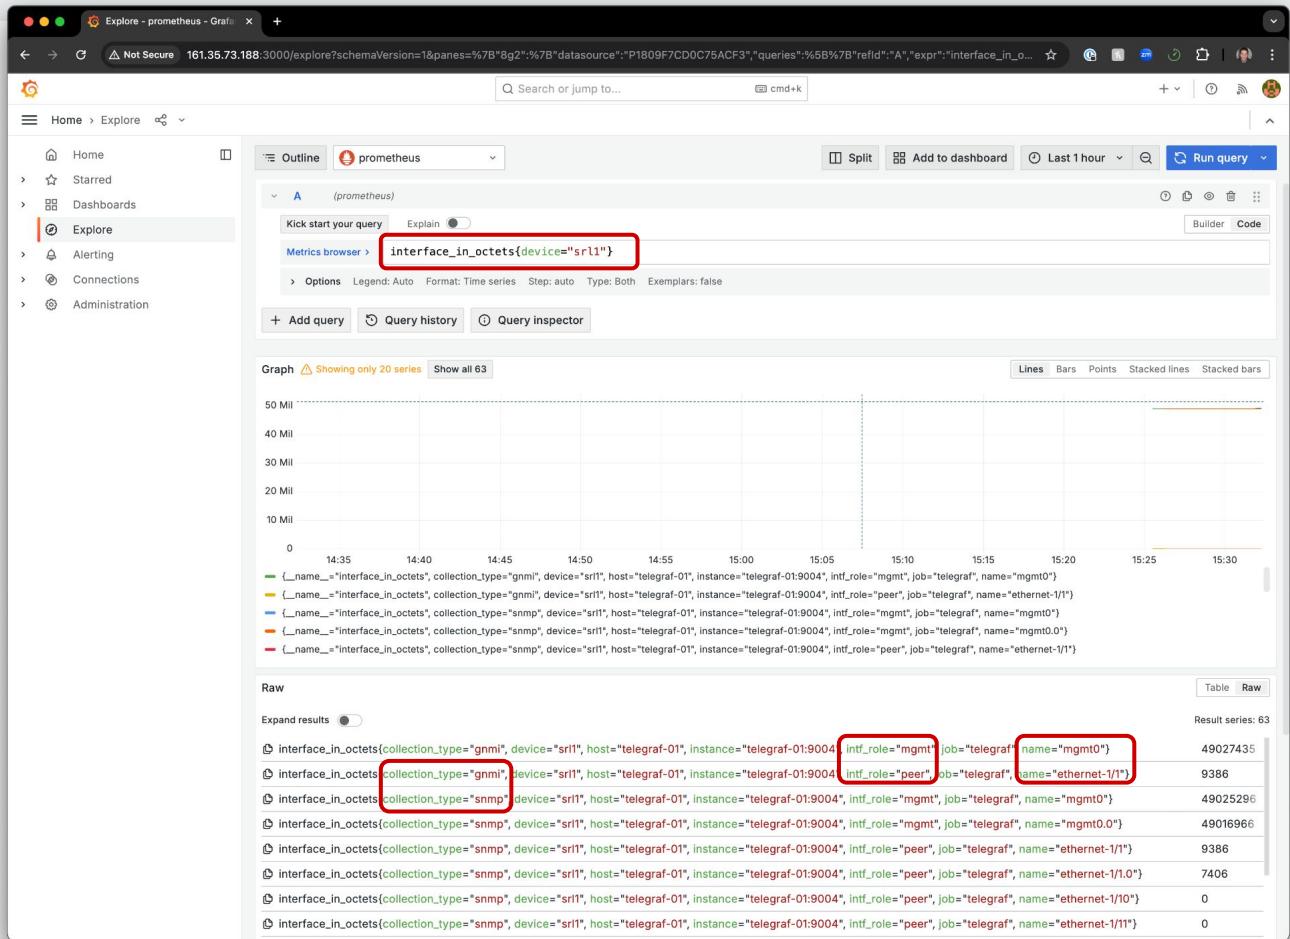
Getting Telegraf outputs (processed)

```
root@netobs-droplet:~/network-observability-lab# netobs docker logs telegraf-01 --tail 10
[18:11:01] Showing logs for service(s): ['telegraf-01']

    Running command: docker compose --project-name netobs -f chapters/batteries-included/docker-compose.yml logs --tail=10 telegraf-01
telegraf-01 | interface,collection_type=gnmi,device=srl1,host=telegraf-01,intf_role=peer,name=ethernet-1/58 oper_state=2i,admin_state=2i 1768143286055505407
telegraf-01 | interface,collection_type=gnmi,device=srl1,host=telegraf-01,intf_role=mgmt,name=mgmt0 oper_state=1i,admin_state=1i 1768143286055505407
telegraf-01 | cpu,collection_type=gnmi,device=srl1,host=telegraf-01,index=all,slot=A used=41i 1768143286057268129
telegraf-01 | memory,collection_type=gnmi,device=srl1,host=telegraf-01,slot=A utilization=67i 1768143286057436961
telegraf-01 | interface,collection_type=snmp,device=srl1,host=telegraf-01,intf_role=peer,name=ethernet-1/50 in_octets=0i,out_octets=0i 17681425800000000000
telegraf-01 | interface,collection_type=snmp,device=srl1,host=telegraf-01,intf_role=peer,name=ethernet-1/30 out_octets=0i,in_octets=0i 17681425800000000000
telegraf-01 | interface,collection_type=snmp,device=srl1,host=telegraf-01,intf_role=peer,name=ethernet-1/32 in_octets=0i,out_octets=0i 17681425800000000000
telegraf-01 | ping,collection_type=ping,device=srl1,host=telegraf-01,service=srl1,watcher_type=healthcheck
minimum_response_ms=4.032,maximum_response_ms=4.594,standard_deviation_ms=0.231,percent_packet_loss=0,ttl=64i,average_response_ms=4.336,result_code=0i,packets_transmitted=3i,packets_received=3i 1768142592000000000
```

End of task: show logs

Verify Metrics in Storage



Exercise 2.2: Build Your First Normalized Telemetry Pipeline Logs with Logstash

Logstash Input

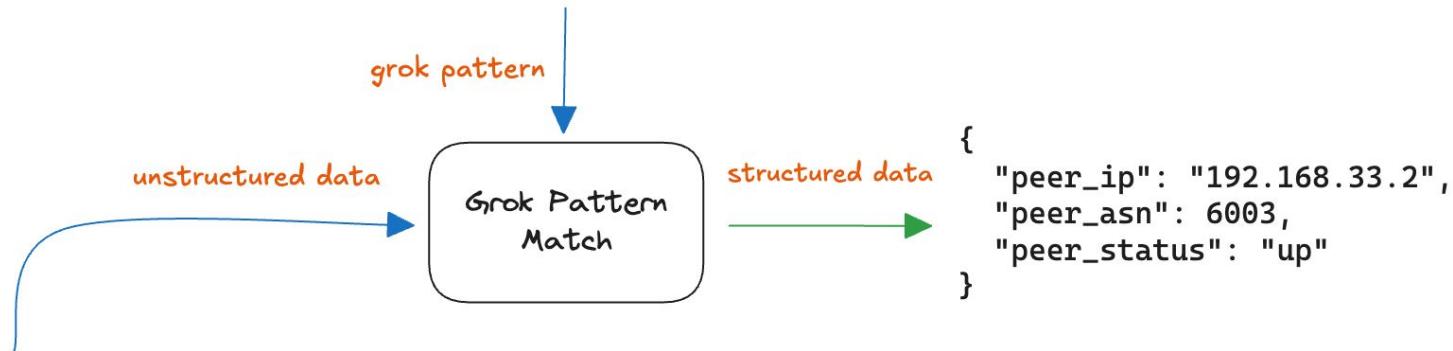
```
input {
    syslog {
        port => 1515
        use_labels => true
        id => "srlinux"
        tags => [ "syslog", "srlinux" ]
        timezone => "Europe/Rome"
    }
}
```

Other options:

- beats
- raw tcp/udp
- http
- kafka
- file

Logstash Grok

```
peer %{IPV4:peer_ip}, AS %{INT:peer_asn}, is %{WORD:peer_status}
```



Logstash Grok and Mutate

```
filter {
  if [type] == "syslog" {
    grok {
      patterns_dir => [ "/var/lib/logstash/patterns" ]
      match => { "message" =>
        "%{SRLPROC:subsystem}\|%{SRLPID:pid}\|%{SRLTHR:thread}\|%{SRSLSEQ:sequence}\|%{SRLLVL:initial}:\s+(?<log_message>(.*))"
      }
    mutate {
      rename => { "severity_label" => "level" }
      lowercase => [ "level" ]
    }
    # Add a field for metadata example
    mutate {
      add_field => {
        "netpanda" => "bethepacket"
      }
    }
  }
}
```

Other options:

- json
- date
- geoip
- drop
- translate
- ruby

Logstash Output

Other options:

- elasticsearch
- kafka
- stdout
- filte
- http

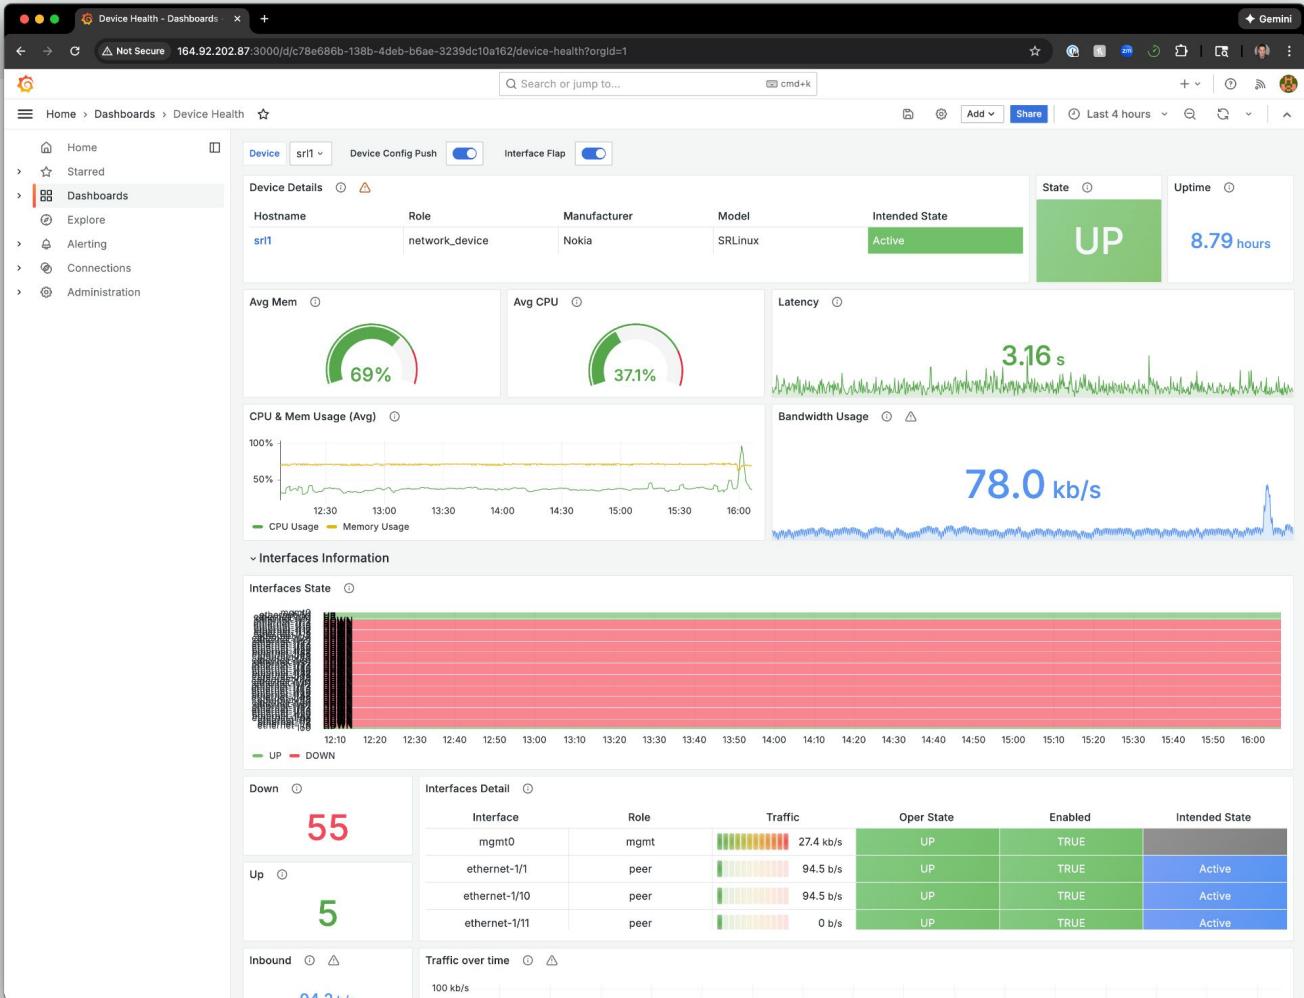
```
output {
    if [type] == "syslog" {
        loki {
            url => "http://loki:3001/loki/api/v1/push"
            metadata_fields => ["netpanda"]
            include_fields => [
                "device",
                "level",
                "facility",
                "facility_label",
                "vendor_facility",
                "vendor_facility_process",
                "event_type",
                "interface",
                "interface_status",
                "priority",
                "severity",
                "type"
            ]
        }
    }
}
```

Getting Logstash outputs

```
root@netobs-droplet:~/network-observability-lab# netobs docker logs logstash --tail 16
[16:52:40] Showing logs for service(s): ['logstash']
    Running command: docker compose --project-name netobs -f
    chapters/batteries-included/docker-compose.yml logs --tail=16 logstash
    logstash | {
    logstash |     "message" => "<180>1 2026-02-01T16:52:27.984497+00:00 srl1
    sr_xdp_lc_1 --- debug|2639|3138|00170|W: csim_pd  csim_platform.cc:2553
    UpdateLicenseValidity No valid license, limiting packet rate to 10000pps\n",
    logstash |     "facility" => 0,
    logstash |     "host" => "198.51.100.2",
    logstash |     "severity" => 0,
    logstash |     "priority" => 0,
    logstash |     "@version" => "1",
    logstash |     "@timestamp" => 2026-02-01T16:52:27.989Z,
    logstash |     "facility_label" => "kernel",
    logstash |     "tags" => [
    logstash |         [0] "syslog",
    logstash |         [1] "srlinux",
    logstash |         [2] "_grokparsefailure_sysloginput"
    logstash |     ],
    logstash |     "severity_label" => "Emergency"
    logstash | }
```

Successfully ran: show logs

Visualize Metrics and Logs in a Graph





Wrap Up

In this scenario:

- You learn about the challenges of network data collection
- The basic collector architecture is INPUT -> PROCESS -> OUTPUT
- You implemented metrics collection with Telegraf and logs with Logstash
 - Pick the right tool for each occasion
- Data collection is just the beginning of the Observability journey

3

Turning Data Into Insight: Alerting, Queries & Visualization

Interpretation layer: queries → rules → alerts → dashboards



Turning Data Into Insight: Alerting, Queries & Visualization

- In Section 2 we *collected* metrics + logs
- In this sections we will learn how to **use** them:
 - Explore the raw signals (metrics + logs)
 - Query them into answers
 - Alert on decisions (intent vs reality)
 - Visualize health + history + context in dashboard

Grafana - <http://<lab-address>:3000>

User: **netobs**

Password: **netobs123**

Exploring the data: Metrics taxonomy

- Metric name: `interface_oper_state`
- Labels:

- `device="srl1"` (which box/device)
- `name="ethernet-1/1"` (which interface)
- `intf_role="peer"` (semantic role from enrichment)
- Sample value: 2 (meaning depends on mapping)
- Timestamp: (implicit per-sample)

Queries interpret signals into answers:

How many interfaces are **up** per device?

```
count by (device) (
    interface_oper_state{intf_role="peer"} == 1
)
```



`interface_oper_state {device="srl1", name="ethernet-1/1", intf_role="peer"} 2`

The diagram shows the components of a metric definition. It consists of three parts: "Metric Name" (the metric name itself), "Labels" (the labels attached to the metric), and "Value" (the sample value). A black bracket groups the "Labels" and "Value" parts. A green number "2" is placed next to the "Value" part. A black arrow points from the "Metric Name" to the "Labels" group. Another black arrow points from the "Value" part to the green "2".

1 = UP
2 = DOWN

Exploring the data: Prep - traffic generation

```
> ssh admin@srl11
Password: NokiaSrl11!
Welcome to the Nokia SR Linux CLI.

--{ + running }--[ ]--

A:admin@srl11#
--{ + running }--[ ]--

A:admin@srl11# ping 10.1.2.2 network-instance default -s 1400 -M do -i
0.5
```

Exploring the data: PromQL

Query: What is the inbound throughput of a network device?

1. We identify the metric (`interface_in_octets`) and filter for an specific device and interface

```
interface_in_octets{device="srl1", name="ethernet-1/1"}
```

2. Metrics are of type **counter**, and we need to calculate the traffic using `rate()` function

```
rate(interface_in_octets{device="srl1", name="ethernet-1/1", collection_type="gnmi"}) [2m] * 8
```

3. Now we can `sum()` the traffic of all in-band interfaces to get the inbound throughput of a device

```
sum(rate(interface_in_octets{device="srl1", name!~"mgmt0.*", collection_type="gnmi"} [2m])) * 8
```

4. Now per device

```
sum by (device) rate(interface_in_octets{name!~"mgmt0.*", collection_type="gnmi"} [2m]) * 8
```

Exploring the data: PromQL (Part 2)

Query: What are the interfaces of the devices that are enabled but operationally up?

1. Check **interfaces state is enabled** (these are enum-like metrics where the value tells you the state)

```
interface_admin_state{intf_role="peer"} == 1
```

2. Check **interfaces operational state = up**

```
interface_oper_state{intf_role="peer"} == 1
```

3. Match the interfaces and the devices from both views

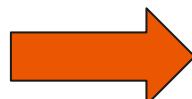
```
and on (device, name)
```

4. Group results by device + interface

```
count by (device, name) ( ... )
```

Resulting Query:

```
count by (device, name) (  
    interface_admin_state{intf_role="peer"} == 1  
    and on (device, name)  
    interface_oper_state{intf_role="peer"} == 1  
)
```



Exercise 3.1: Explore the Metrics Data

Lab 3.1 Metrics Discovery

1. "How many peer interfaces are ENABLED vs DISABLED on a device?"
2. "How many BGP sessions are enabled but operationally down?"
 - a. HINT: bgp_admin_state and bgp_oper_state
 - b. HINT: Look for the value-mapping in your telegraf configuration)
3. "What is the total amount of BGP IPv4 received routes from a neighboring device?"
 - a. HINT: bgp_received_routes
 - b. HINT: Filter on afi_safi_name="ipv4-unicast"
4. "Show the interfaces that have flapped more than 3 times in the span of 2 minutes"
 - a. HINT: interface_carrier_transitions is a counter for interface state changes on its carrier.
 - b. HINT: Use rate/increase functions to determine the rate of change in the span of 2 minutes
 - c. HINT: Trigger interfaces flap with `-netobs utils device-interface-flap --device srl1 --interface ethernet-1/1 --count 10 --platform nokia_srl`

Exploring the data: Logs taxonomy

- A log entry is defined by labels (indexed)

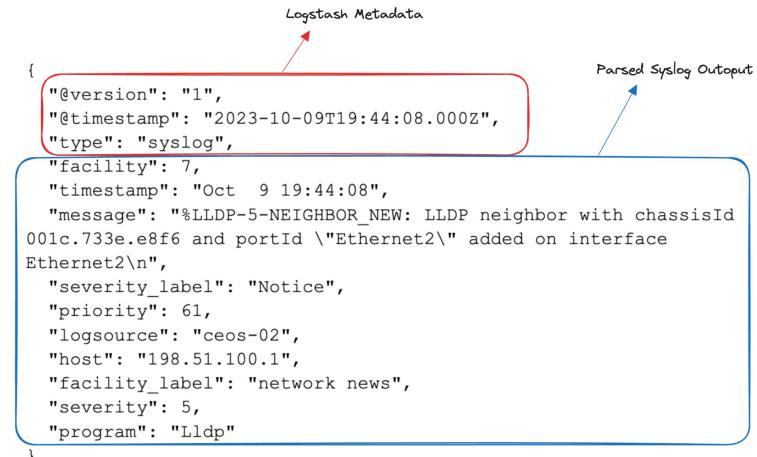
```
{device="sr11", job="netdev-logs", facility_label="srlinux"}
```

- A log content is the message/line content (not indexed)
- Rule of thumb:
 - Labels = fast stream selection (indexed)
 - Line fields/message = filter after selection
- LogQL taxonomy



Stream Selector

```
{ label1="value1", label2="value2" } <filter operator> <log pipeline expression>
```



Log Pipeline (Optional)

Exploring the data: Prep - interface flap

```
# Shut/Unshut interface 10 times with a delay of 10 seconds

> netobs utils device-interface-flap \
    --device srl1 \
    --interface ethernet-1/1 \
    --count 10 \
    --delay 10 \
    --platform nokia_srl
```

Exploring the data: LogQL

Query: How frequent do we have someone enabling/disabling interfaces?

1. We identify the logs for **srl1**

```
{device="srl1"}
```

2. Filter license-related messages to reduce noise

```
{device="srl1"} != "license"
```

3. Next we are interested in the **admin-state** changes that we have made on the device

```
{device="srl1"} != "license" |= "set / interface" |= "admin-state"
```

4. Now, let's create a metric of the rate of changes performed on the interfaces of the device

```
sum(rate({device="srl1"} != "license" |= "set / interface" |= "admin-state" [5m]))
```

Exercise 3.2: Explore the Logs

Data

Lab 3.2 Logs Discovery

1. “Can I get only the **enable** events from the interfaces?”
2. “Can you filter only when an user is connected to the device and running commands?”
 - a. HINT: **aaa** and **debug** modules provide overall view of commands and VTY sessions to the device

Actioning on data: Alerts

- A **query** answers a question
- An **alert** is the same query plus: “if true -> act”

Decision: “Alert when a peer interface is configured UP but operationally DOWN”

```
count by (device, name) (
  interface_admin_state{intf_role="peer"} == 1)
  and on (device, name)
  interface_oper_state{intf_role="peer"} == 2
) > 0
```



- Intent: admin_state == 1 (we expect it UP)
- Reality: oper_state == 2 (it's actually DOWN)
- and on(device, name) ensures we compare the same interface

Actioning on data: Alerts (cont.)

With decision at hand - we create a rule:

- Evaluate **for: 2m** to avoid flappy noise
- Add **annotations** so humans know what to do.
- Add **labels** to provide more context or categorization for later processing.

```
groups:
- name: interface_intent_mismatch
rules:
- alert: InterfaceAdminUpOperDown
  expr: |
    count by (device, name) (
      interface_admin_state{intf_role="peer"} == 1
      and on (device, name)
      (interface_oper_state{intf_role="peer"}) == 2
    ) > 0
  for: 2m
  labels:
    severity: warning
    category: network
  annotations:
    summary: "Interface intent mismatch detected"
    description: |
      Interface {{ $labels.name }} on device {{ $labels.device }}
      is configured as UP (admin state) but is currently DOWN (oper state).

      This usually indicates a cabling, peer, or physical-layer issue.
```

Exercise 3.3: Actioning on Data

- Alerts



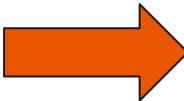
Lab 3.3: Alerts

1. Activate the alert rule
2. Update Prometheus rules configuration:
 - a. HINT: You can use Prometheus reload API endpoint or by running the command: **netobs lab update prometheus**
3. Check alert being triggered in Prometheus <http://<lab-address>:9090/#/alerts>
4. After alert is firing check Alertmanager <http://<lab-address>:9093/>

Actioning on data: Create Visualizations

- Use queries to tell the **operational story**
- Ask - which story are we telling?

For one device, we want to answer:

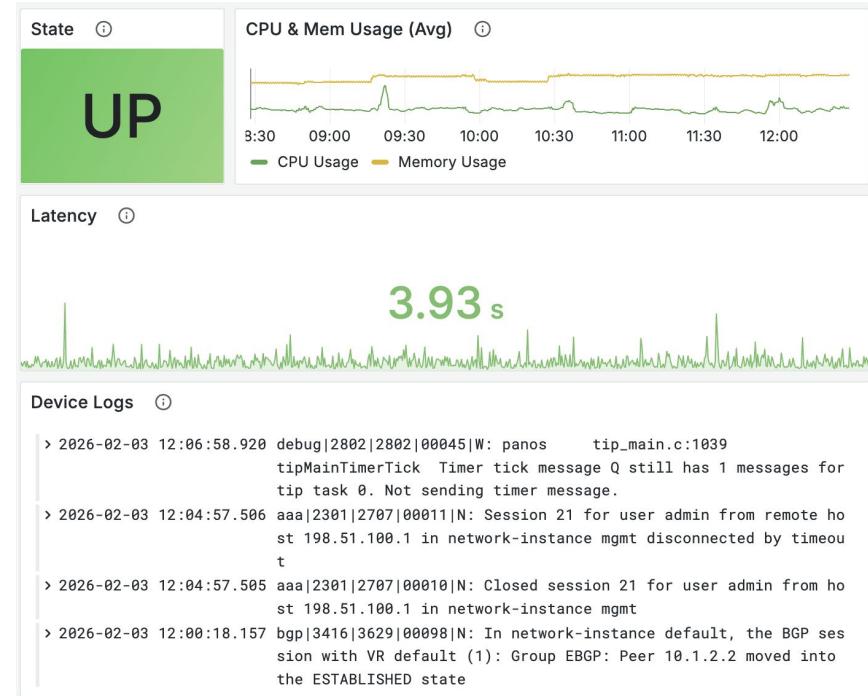


- How healthy is it right now? (summary)
- When did it change? (timeline)
- What happened around it? (logs/context)

NOTE: Grafana - <http://<lab-address>:3000>

User: **netobs**

Password: **netobs123**



Actioning on Data: A Story of Interfaces

Let's create a dashboard for telling the story of the interfaces in a device operationally.

- Dashboard Settings > Title
- Variables > New Variable
 - Variable Type: Query
 - Name: `device`
 - Label: `Device`
 - Query Options:
 - Query Type: `Lavel Values`
 - Label*: `device`
 - Metric: `device_uptime`

Actioning on Data: Interfaces Admin Panel

Visualization:

- Type: State Timeline
- Title: Interfaces Admin State
- Value Mapping:
 - 1 -> Enabled -> Blue
 - 2 -> Disabled -> Purple
- Standard Options > Color Scheme: Single Color

Query:

- Datasource: prometheus
- Legend: {{ name }}

```
interface_admin_state{  
    intf_role="peer",  
    name=~"ethernet-1/1([0-2])?",  
    device="$device"  
}
```

Exercise 3.4: Actioning on Data

- Dashboards



Lab 3.4: More Interfaces Panels...

Create an Interface Logs Panel

- Type: Logs
- Datasource: Loki
- Title: Interface Logs -\$device

```
{device="$device" } != "license"  
|= "interface"  
|~ "down|up|admin-state"  
| line_format "{{.device}} {{.facility_label}} | {{_line_}}"
```

Lab 3.4: More Interfaces Panels...

Create an Interface Oper State Panel

- Type: State Timeline
- Datasource: prometheus
- Title: Interface Operating State -\$device
- Value Mapping:
 - 1 -> UP -> Green
 - 2 -> DOWN -> Red
- Legend: {{name}}

```
interface_oper_state{  
    intf_role="peer",  
    name=~"ethernet-1/1([0-2])?",  
    device="$device"  
}
```

Lab 3.4: More Interfaces Panels...

Create an Interface Traffic Panel

- Type: Timeseries
- Datasource: prometheus
- Title: Interfaces Traffic
- Visibility:
 - Mode > Table
 - Placement > Right
 - Values > Last*
- Standard Options > Units: bits/sec (SI)
- Legend: In|Out: {{name}}

```
# Series A: IN Traffic
rate(
    interface_out_octets{name=~"ethernet-1/1([0-2])?", device="$device", collection_type="gnmi"}[$__rate_interval]
) * 8

# Series B: OUT Traffic
rate(
    interface_in_octets{name=~"ethernet-1/1([0-2])?", device="$device", collection_type="gnmi"}[$__rate_interval]
) * 8 * -1
```

4

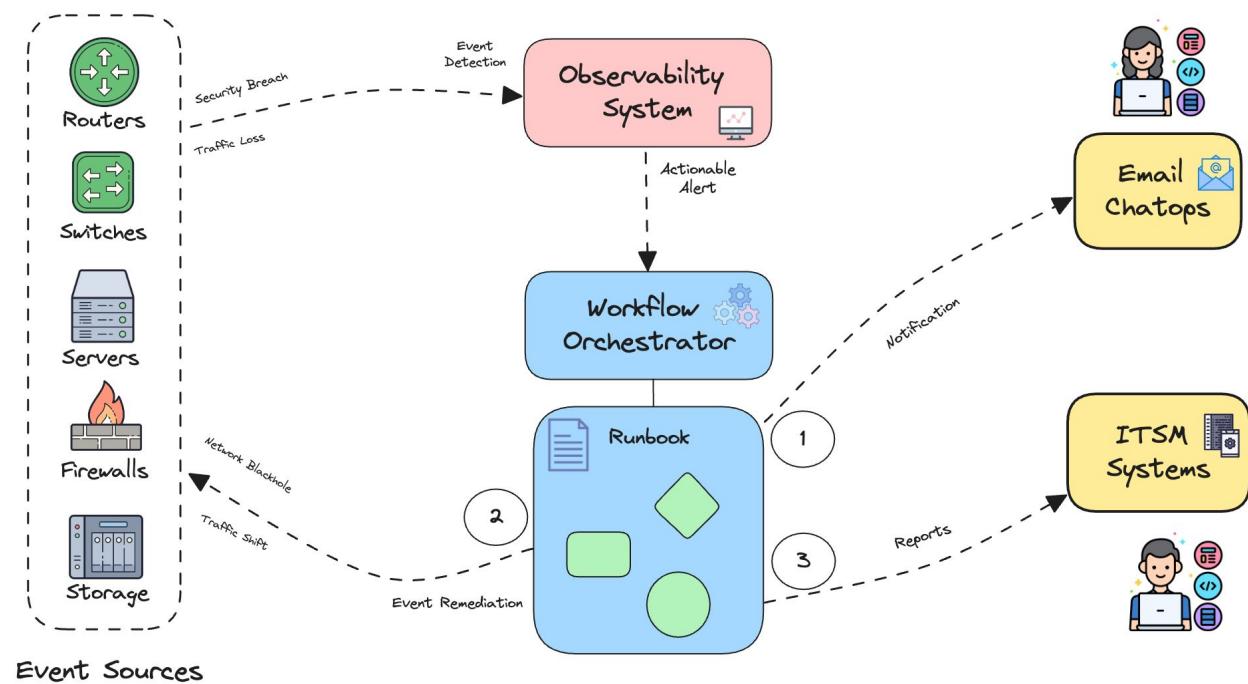
Day-2 Automation: From Observability to Intelligent Workflows

Day-2 Automation starts with *evidence*, not actions

- Automation is a **decision loop**, not “runbook as code”.
- The workflow must be **safe-by-default** (gates, context, audit trail).

The standard view for day-2:

- Monitoring pages you.
- Observability gives you enough information to act.
- **Automation is the “Day-2” layer** sitting on top of observability data.

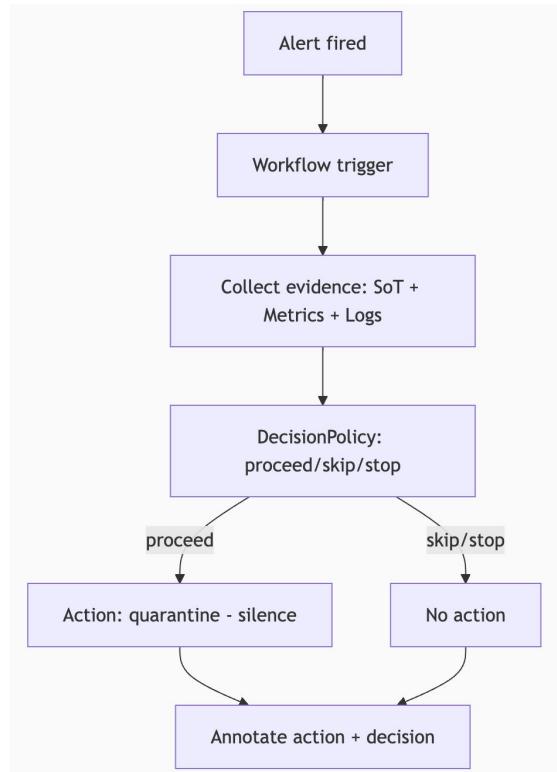


Observability -> Automation Loop

- Alerts are triggers
- Evidence is collected
- Policy decides
- Action + annotation closes the loop.

Workshop Flow:

1. Explore Workshop SDK and its features
2. Play around with the methods and data
3. Move to Prefect workflows development
4. Trigger and validate executions



Auto: Prep - Prefect Server

```
# Navigate to webinar folder
cd ~/network-observability-lab/chapters/webinar

# Setup prefect to talk to the docker prefect server
prefect config set PREFECT_API_URL=http://localhost:4200/api

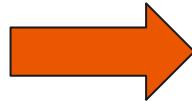
# Check the variable is set
prefect config view
```

Auto: IPython = your observability REPL

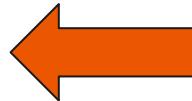
- IPython used to explore SDK and data
- Imports pasted for easier load and prettier outputs

```
# Explore Python SDK with ipython  
  
ipython
```

```
sdk.prom # Prometheus client  
  
sdk.loki # Loki client  
  
sdk.nb # Nautobot client  
  
sdk.am # Alertmanager client
```



```
from rich.pretty import install  
  
install()  
  
  
from netobs_workshop_sdk import (  
  
    WorkshopSDK,  
  
    EvidenceBundle,  
  
    DecisionPolicy,  
  
    Decision,  
)  
  
sdk = WorkshopSDK()
```



Auto: Start with intent

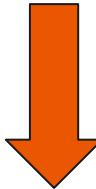
“What should be true” **must gate** “what we do” —> SoT + maintenance + peer intent gate

```
# Get a Device object from SoT  
  
dev = sdk.nb.get_device("srl1")  
  
dev["name"]
```



```
# Create SoT gate  
  
gate = sdk.nb.build_bgp_intent_gate(  
    device="srl1",  
    peer_address="10.1.2.2",  
    afi_safi="ipv4-unicast",  
)  
  
gate
```

```
# Get a Device object from SoT  
  
sdk.nb.is_device_in_maintenance(dev)
```

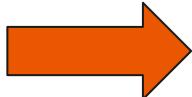


Auto: Reality check - what does the metrics say?

With PromQL you can see and the **metrics snapshot** you can reason on

```
# See PromQL queries

qs = sdk.bgp_queries(
    device="srl1",
    peer_address="10.1.2.2",
    afi_safi="ipv4-unicast",
    instance_name="default",
)
qs
```



```
# Run queries

# Check raw results
for k, q in qs.items():
    res = sdk.prom.instant(q)
    print(k, "=>", res[:1])
```

```
# Get a structured data snapshot

sdk.bgp_metrics_snapshot(
    device="srl1",
    peer_address="10.1.2.2",
    afi_safi="ipv4-unicast",
    instance_name="default",
)
```

Auto: Reality check - what does the logs add?

Metrics tell you **what**; logs often explain **why**.

```
# If needed generate some logs  
  
cd ~/network-observability-lab  
  
# Run the flap interface command  
  
netobs utils device-interface-flap \  
    --device srl1 \  
    --interface ethernet-1/1 \  
    --count 2 \  
    --delay 2 \  
    --platform nokia_srl
```

```
# Check the LogQL used  
  
sdk.bgp_logql(device="srl1", peer_address="10.1.2.2")
```



```
# Check some logs  
  
logs = sdk.bgp_logs(  
    device="srl1", peer_address="10.1.2.2", minutes=30, limit=50  
)  
  
len(logs), logs[:2]
```

Auto: Correlation -> build an EvidenceBundle

Correlation is where observability becomes automation-grade.

```
# Collect BGP Evidence

ev = sdk.collect_bgp_evidence(
    device="srl1",
    peer_address="10.1.2.2",
    afi_safi="ipv4-unicast",
    instance_name="default",
    log_minutes=30,
    log_limit=50,
)
ev.summary()
```



```
# Check the raw data collected

ev.sot
ev.metrics
ev.logs[:3]
```

Auto: DecisionPolicy -> “encoded decisions”

Policies exist to make “why” explicit and auditable.

```
# SoT only gate  
  
policy = DecisionPolicy()  
  
policy.evaluate(ev.sot)
```



```
# Intent vs Reality  
  
policy = DecisionPolicy()  
  
policy.evaluate(ev.sot, ev.metrics)
```

```
# Try with a broken peer  
  
ev_broken = sdk.collect_bgp_evidence(  
    device="srl2",  
    peer_address="10.1.11.1",  
    afi_safi="ipv4-unicast",  
    instance_name="default",  
    log_minutes=30,  
    log_limit=50,  
)  
  
DecisionPolicy().evaluate(ev_broken.sot, ev_broken.metrics)
```

Action: quarantine via Alertmanager silence

Actions are the last step, not the first → and only after a proceed

```
# Quarantine action - silence alert

silence_id = sdk.quarantine_bgp(
    device=ev.device, peer_address=ev.peer_address, minutes=20
)

silence_id
```

Check Silences in <http://<lab-address>:9093/#/silences>

Close the loop: annotate decisions + actions

If you can't explain it later, it's unsafe automation.

```
# General annotation

sdk.annotate(
    labels={
        "source": "prefect",
        "workflow": "demo_quarantine_bgp",
        "device": ev.device,
        "peer_address": ev.peer_address,
    },
    message=f"QUARANTINE applied (silence_id={silence_id})",
)
```

```
# Decision-specific annotation

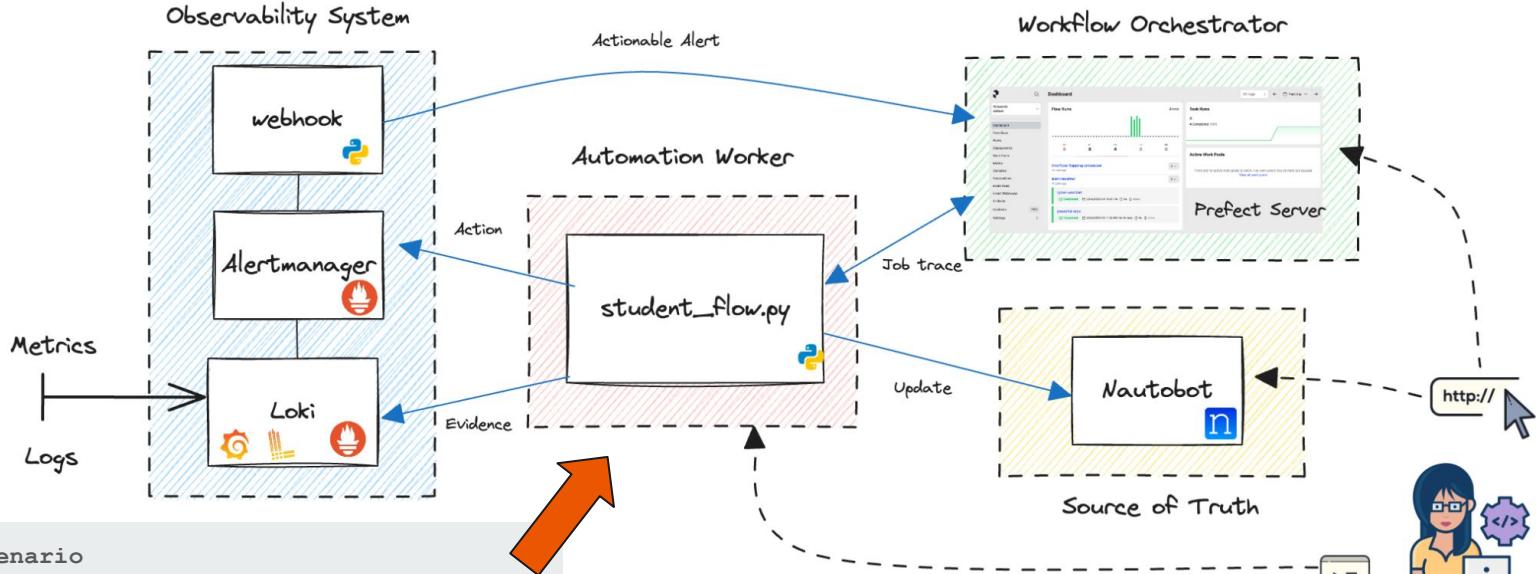
sdk.annotate_decision(
    workflow="demo_quarantine_bgp",
    device=ev.device,
    peer_address=ev.peer_address,
    decision="skip",
    message="SKIP quarantine: peer not intended in SoT",
)
```

LLM-ready RCA payload (structured + scoped)

RCA becomes feasible when evidence is already correlated + normalized.

```
# Payload for RCA  
  
payload = ev.to_rca_payload()  
  
payload.keys()
```

Run it on Prefect



```
# Navigate to the scenario
cd ~/network-observability-lab/chapters/webinar-completed

# Prepare the runner

python student_flow.py
```

Trigger 1: Actionable (expect quarantine)

```
prefect deployment run 'alert-receiver/alert-receiver' \  
    --param alertname='BgpSessionNotUp' \  
    --param status='firing' \  
    --param alert_group='{"status":"firing","groupLabels":{"alertname":"BgpSessionNotUp"},  
    "alerts":[{"labels":{"alertname":"BgpSessionNotUp","device":"srl2","peer_address":"10.1.11.1",  
    "afi_safi_name":"ipv4-unicast","name":"default"},  
    "annotations":{"summary":"expect QUARANTINE (mismatch demo)"}}]}'
```

Trigger 2: Skip (healthy peer)

```
prefect deployment run 'alert-receiver/alert-receiver' \
    --param alertname='BgpSessionNotUp' \
    --param status='firing' \
    --param alert_group='{"status":"firing","groupLabels":{"alertname":"BgpSessionNotUp"}, \
    "alerts":[{"labels":{"alertname":"BgpSessionNotUp","device":"srl2","peer_address":"10.1.2.1", \
    "afi_safi_name":"ipv4-unicast","name":"default"}, \
    "annotations":{"summary":"expect SKIP (peer is healthy)"}}]}'
```

Trigger 3: Resolved Path (alert is resolved)

```
prefect deployment run 'alert-receiver/alert-receiver' \
    --param alertname='BgpSessionNotUp' \
    --param status='resolved' \
    --param alert_group='{"status":"resolved","groupLabels":{"alertname":"BgpSessionNotUp"}, \
"alerts":[{"labels":{"alertname":"BgpSessionNotUp","device":"srl2","peer_address":"10.1.11.1", \
"afi_safi_name":"ipv4-unicast","name":"default"}, \
"annotations":{"summary":"resolved test"}}]}'
```

Exercise 4: Actioning on Data - Event Driven Automation

Lab 4: Event Driven Automation Workflows

1. Work with the lab environment to create an alert that triggers an automation workflow that:
 - a. is actionable (quarantine)
 - b. is skippable
 - c. is resolved
2. What happens if a device is in `maintenance` mode from the SoT?
 - a. Change a device into maintenance mode and verify the logic follows the pattern expected



Thank you!